

More on Pointers, Aggregate types, Files

COMP2129

FACULTY OF
ENGINEERING &
INFORMATION
TECHNOLOGIES



THE UNIVERSITY OF
SYDNEY



- › Revision pointers
 - › Misc C operators and types
 - › Aggregate types
 - struct
 - › Files
 - Finally!
-

› What is the value held by p? and how much memory is used by p (in bytes)?

› `int p;`

› `char p;`

› `void foo(int *p)`

› `char *p;`

› `char **p;`

› What is the value held by p? and how much memory is used by p (in bytes)?

› `int p;`

› `char p;`

› `void foo(int *p)`

› `char *p;`

› `char **p;`

› `int **p;`

› `long *p;`

› `void *p;`

› `const unsigned long long int * const p;`

› `bubblebobble *****p;`

› `char *p`

- Address to a single char value
- Address to a single char value that is the first in an array

› `char *argv[]`

- Array of “the type” with unknown length
- Type is `char *`

› `char **argv`

- `*` Address to the first element to an array of type `char *`
- Then, each element in `*` is an...
 - `**` address to the first element to an array of type `char`

- › Interpretations of `int **data;`
 1. Pointer to pointer to single int value
 2. Array of addresses that point to a single int
 3. Address that points to one array of int values
 4. Array of addresses that point to arrays of int values
-

- › Interpretations of `int **data;`
 1. Pointer to pointer to single int value
 2. Array of addresses that point to a single int
 3. Address that points to one array of int values
 4. Array of addresses that point to arrays of int values

 - › Thinking about each `*` as an array:
 1. Array size ==1, Array size ==1
 2. Array size >=1, Array size == 1
 3. Array size ==1, Array size >= 1
 4. Array size >=1, Array size >= 1
-

- › When you call a function in Java, compare passing a primitive type and Object type.
- › You may have heard:
 - Pass by value
 - Pass by reference

What is the meaning of this in C?

- › void has no size, but `sizeof(void*)` is the size of an address
 - › Pointers are unsigned numbers, why?
-




Pointer arithmetic

› `int *p = NULL;`

› `int x[4];`

› `p = x;`

`x = 0x0100`




› Seeking to the nth byte from a starting address?

› `int *p = NULL;`

› `int x[4];`

› `p = x;`

`x = 0x0100`



› Seeking to the nth byte from a starting address?

```
void *get_address( sometype *data , int n) {  
    unsigned char *ptr = (unsigned char*)data;  
    return (void*) (ptr + n);  
}
```

- › Not all h/w architectures are the same
 - different sizes for basic types
 - › C specification does not dictate exactly how many bytes an int will be
 - › **sizeof** operator returns the number of bytes used to represent the given type or expression
 - sizeof(char)
 - sizeof(int)
 - sizeof(float *)
 - sizeof (1)
 - sizeof(p)
-

- › Not all h/w architectures are the same
 - different sizes for basic types
 - › C specification does not dictate exactly how many bytes an int will be
 - › **sizeof** operator returns the number of bytes used to represent the given type or expression.
 - **sizeof**(**char**)
 - **sizeof**(**int**), **sizeof**(**double**)
 - **sizeof**(**float** *)
 - **sizeof** (1), **sizeof** (1/2), **sizeof** (1.0 / 2.0)
 - **sizeof**(p) ????
-

› Special case for **p**, what is it?

- **char** p;
- **char** *p;
- **char** p[8];

› But...

- **char** msg[100];
- **char** *p = msg;
- **char** msg2[] = "got my goat and up my nose";
- **char** *p = msg2;
- **char** *p = "im finna spaz on anybody now"

› **sizeof** needs to be used carefully

- › The types `char` will support the value range from `CHAR_MIN` to `CHAR_MAX` as defined in file `<limits.h>`

```
- #define UCHAR_MAX      255          /* max value for an unsigned char */  
- #define CHAR_MAX       127          /* max value for a char */  
- #define CHAR_MIN       (-128)      /* min value for a char */
```

- › Most C implementations default types as `signed` values, but a warning that you should not assume this.
 - › `unsigned` and `signed` enforce the sign usage
 - `char ch;`
 - `signed char ch;`
 - `unsigned char ch;`
 - `unsigned int total;`
-

- › `const` prevents the value being modified
 - `const char *fileheader = "P1"`
 - `fileheader[1] = '3';` **Illegal: change of char value**

 - › It can be used to *help* avoid arbitrary changes to memory

 - › The value `const` protects depends where it appears
 - `char * const fileheader = "P1"`
 - `fileheader = "P3";` **Illegal: change of address value**

 - › Reading right to left:
 - Is an address, points to a char, that is constant
 - Is an address, that is constant
-

- › `const` prevents the value being modified
 - `const char *fileheader = "P1"`
 - `fileheader[1] = '3';` **Illegal: change of char value**
- › It can be used to *help* avoid arbitrary changes to memory
- › The value `const` protects depends where it appears
 - `char * const fileheader = "P1"`
 - `fileheader = "P3";` **Illegal: change of address value**
- › You can **cast** if you know if the memory is writable

Non-writable

```
char fileheader[] = {'P', '1'};
const char *dataptr = (char*)fileheader;
char *p = (char*)dataptr;
p[1] = '3';
```

writable

- › Exact bit representation unknown, usually IEEE 754
- › Generally, floating point number x is defined as:

$$x = sb^e \sum_{k=1}^p f_k b^{-k}, \quad e_{\min} \leq e \leq e_{\max}$$

- › s sign
- › b base of exponent (e.g. 2, 10, 16)
- › e exponent
- › p precision
- › f_k nonnegative integer less than b

+0

-0

+ve / 0 = +infinite

-ve / 0 = -infinite

NaN (not a number)

Zero exponents...

Enums

COMP2129

FACULTY OF
ENGINEERING &
INFORMATION
TECHNOLOGIES



THE UNIVERSITY OF
SYDNEY



- › simple data types:
 - int, char, float.....

 - › pointers to simple data types:
 - int *, char *, float *
-

- › enums (enumerated types) are another simple type
 - › enums map to int
 - › an enum associates a name with a value
-

```
enum day_name  
{  
    Sun, Mon, Tue, Wed, Thu, Fri, Sat, day_undef  
};
```

- › Maps to integers, 0 .. 7
 - › Can do things like 'Sun ++'
 - › very close to int
-



```
enum month_name  
{  
    Jan, Feb, Mar, Apr, May, Jun,  
    Jul, Aug, Sep, Oct, Nov, Dec,  
    month_undef  
};
```

- › we could always use integers to represent a set of elements
 - › but enums make your code much more readable
 - › eg red instead of 0
 - How many bytes for an array of enum?
-

Structures

COMP2129

FACULTY OF
ENGINEERING &
INFORMATION
TECHNOLOGIES



THE UNIVERSITY OF
SYDNEY



- › So far the only collection of data we've covered is the *array*
 - › Arrays are used to hold items of the **same type** and access them by giving an index
 - › Sometimes we want to hold a collection of data items of ***different*** types.
 - › For example: a library catalogue for a book might contain the title, author's name, call number, date acquired, date due back etc
 - › For this type of collection C has a data type called a ***structure***
-

Structure definition example

struct date

{

enum day_name day;

int day_num;

enum month_name month;

int year;

};

name of the type of structure

fields of the structure

```
struct date {
    enum day_name    day;
    int              day_num;
    enum month_name  month;
    int              year;
} Big_day {
    Mon, 7, Jan, 1980
};

struct date    moonlanding;
struct date    deadline = {day_undef, 1, Jan,
                           2000};
struct date    *completion;
```



```
struct date {  
    enum day_name    day;  
    int              day_num;  
    enum month_name  month;  
    int              year;
```

Structure definition

```
} Big_day
```

Structure declaration

```
{
```

```
    Mon, 7, Jan, 1980
```

Structure initialisation

```
};
```

```
struct date    moonlanding;  
struct date    deadline = {day_undef, 1, Jan, 2000};  
struct date    *completion;
```



```
struct date  moonlanding;
```

```
struct date  deadline =  
                {day_undef, 1, Jan, 2000};
```

```
struct date  *completion;
```



```
struct car_desc
{
    enum car_cols    colour;
    enum car_make    make;
    int               year;
};
```

```
struct [tag]  
{  
    member-declarations  
  
} [identifier-list];
```

- › Once tag is defined, can declare structs with:

```
struct tag    identifier-list;
```

```
struct date bigday;  
int         theyear;
```

```
theyear = bigday.year
```

A dot used to nominate an element of the structure.


```
struct date bigday;  
struct date * mydate;  
int         theyear;
```

```
mydate = &bigday;
```

If a pointer to the structure is used, then the -> operator indicates the element required.

```
theyear = mydate->year
```



typedef

```
typedef struct date{  
    enum day_name    day;  
    int              day_num;  
    enum month_name  month;  
    int              year;  
} Date;
```



typedef

```
typedef struct date{  
    enum day_name    day;  
    int              day_num;  
    enum month_name  month;  
    int              year;  
} Date;
```



```
typedef struct date{
```

```
    enum day_name      day;  
    int                day_num;  
    enum month_name    month;  
    int                year;
```

```
} Date;
```

```
Date Big_day = {Mon, 7, Jan, 1980};
```

```
Date moonlanding;
```

```
Date dopday = {day_undef, 1, Jan, 2000};
```

```
Date *completion;
```

Struct: function arguments, returns

```
struct customer    s1;
struct salesrep    s2;
struct sale transact(struct customer s1, struct salesrep s2);

struct sale transact(struct customer s1,
                    struct salesrep s2)
{
    struct sale s1;
    ...
    return s1;
}
```



- › `stdio.h`
 - › `time.h`
 - › `stat.h`
 - › `pwd.h`
-

```

struct tm
{
    int tm_sec; /* Seconds.      [0-60] */
    int tm_min; /* Minutes.      [0-59] */
    int tm_hour; /* Hours.          [0-23] */
    int tm_mday; /* Day.            [1-31] */
    int tm_mon;  /* Month.          [0-11] */
    int tm_year; /* Year - 1900.    */
    int tm_wday; /* Day of week.    [0-6] */
    int tm_yday; /* Days in year.   [0-365] */
    int tm_isdst; /* DST indicator */
    long int tm_gmtoff; /* Seconds east of UTC. */
    const char *tm_zone; /* Timezone abbreviation. */
};

struct tm * localtime(long *); /* forward decl. */
struct tm * now;

now = localtime(&sometime);
    /* sometime contains time in seconds after
       Jan 1 1970 */

```




```
Hour_now = now->tm_hour;
```

```
printf ("%d/%d/%d\n", now->tm_mday, now->tm_mon,  
        now->tm_year);
```




Memory alignment


```
struct a {  
    int x;  
    short s1, s2;  
    float y;  
    char c1, c2, c3, c4;  
};
```



```
sizeof (struct a) == 16
```

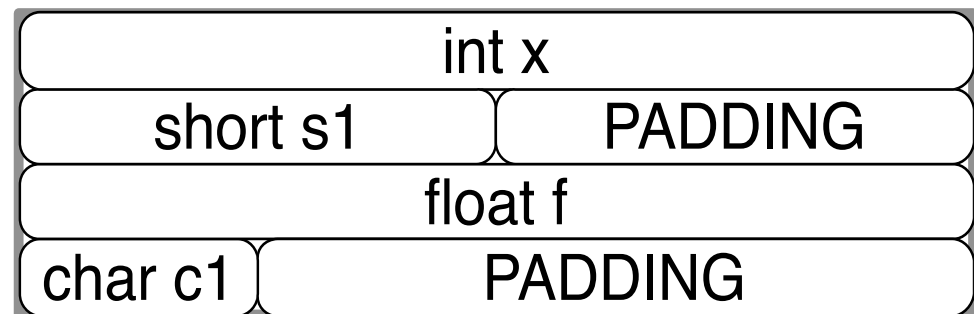
Memory alignment

```
struct a {  
    int x;  
    short s1, s2;  
    float y;  
    char c1, c2, c3, c4;  
};
```



```
sizeof (struct a) == 16
```


```
struct b {  
    int x;  
    short s1;  
    float y;  
    char c1;  
};
```



```
sizeof (struct b) == 16
```


Memory alignment

```
struct b {  
    int x;  
    short s1;  
    float y;  
    char c1;  
};
```



`sizeof (struct b) == 16`


```
struct c {  
    int x;  
    short s1;  
    char c1;  
    float y;  
};
```



`sizeof (struct c) == 12`

- Address of a struct variable will give us direct access to bytes of the first members
 - Alignment depends on architecture
 - Special compiler extensions can be used to prevent padding
 - h/w speed/memory

```
struct c {  
    int x;  
    short s1;  
    char c1;  
    float y;  
};
```



```
sizeof (struct c) == 12
```

Unions

COMP2129

FACULTY OF
ENGINEERING &
INFORMATION
TECHNOLOGIES



THE UNIVERSITY OF
SYDNEY



- › Sometimes we want several variants of a structure but don't want to consume more memory
- › the C *union* lets you declare variables that occupy the **same** memory

- › A library catalogue that contains information about books and films
- › for books we want to store:
 - author
 - ISBN
- › for films we want to store:
 - director
 - producer

```
enum holding_type {book, film};
struct catalog
{
    char * title;
    enum holding_type type;
    struct /* book */
    {
        char * author;
        char * ISBN;
    } book_info;
    struct /* film */
    {
        char * director;
        char * producer;
    } film_info;
};
```

Solution 1

How many bytes
total?

only one of the
structures **book_info**
or **film_info** is used
at any one time.
this can be a major
waste of memory

- › in the first solution, only one of the structures `book_info` or `film_info` is used at any one time.
- › this can be a major **waste of memory**
- › instead, we can use a *union* to indicate that each variant occupies the **same** memory area

```

enum holding_type {book, film};
struct catalog
{
    char * title;
    enum holding_type type;
    union
    {
        struct /* book */
        {
            char * author;
            char * ISBN;
        } book_info;

        struct /* film */
        {
            char * director;
            char * producer;
        } film_info;
    } info;
};

```

Solution 2

we can use a *union* to indicate that each variant occupies the **same** memory area

› to access elements of a union we use the notation
`union_name.part_name`

› example:

union

{

int **a;**

char **b;**

} **x;**

← int →

←char→

11	22	33	44
----	----	----	----

x.a = 0x11223344;

› to access elements of a union we use the notation
`union_name.part_name`

› example:

union

{



int **a;**

char **b;**

} **x;**

← int →

←char→



x.a = 0x11223344;

x.b = 'c';

› in our example, we would access the author this way:

```
struct catalog x;
```

```
x.info.book_info.author
```

- › How can you tell what variant of the union is being used?
- › Answer: you can't!
- › need to have a separate variable to indicate variant in use

Access Example

```
struct catalog x;
```

an enum that indicates the variant

```
switch (x.holding_type)
```

```
{
```

```
    case book:
```

```
        printf("author: %s\n", x.info.book_info.author);
```

```
        break;
```

```
    case film:
```

```
        printf("producer: %s\n", x.info.film_info.producer);
```

```
        break;
```

```
}
```

Files in C

COMP2129

FACULTY OF
ENGINEERING &
INFORMATION
TECHNOLOGIES



THE UNIVERSITY OF
SYDNEY



- › Disk storage peripherals provide persistent storage with a low-level interface
 - Fixed-size blocks
 - Numeric addresses
 - › Operating system arranges this into an **abstraction as files**
 - Files can be variable length
 - Files have names
 - Files have meta-data (owner, last modified date, etc)
 - Files are arranged into eg a tree, by folder/directory structure
 - › Read or write a file is done through System Calls (APIs)
-

- › Devices are often represented as files
 - software reads/write file to access the device
 - E.g. Send a command to the printer by writing to a particular file name
 - › If a file can be a physical device, then it is not fixed in size or behaviour.
 - › A *stream* is associated with a file
 - May support a file position indicator [0, file length]
 - Can be binary or not (e.g. ASCII, multibyte)
 - Can be open/closed/flushed!
 - Can be *unbuffered*, *fully buffered* or *line buffered*
-

- › For each file opened, there needs to be a file descriptor
- › The descriptor describes the state of the file
 - Opened, closed, position etc.
- › `#include <stdio.h>`
 - contains many standard I/O functions and definitions for using files
- › `FILE` is a struct that is defined in `stdio.h` and this is the descriptor
- › To open a file, we use the `fopen` function

```
FILE *fopen(const char *path, const char *mode);
```

```
FILE * myfile = fopen("turtles.txt", "w");
```

- › For each file opened, there needs to be a file descriptor
- › The descriptor describes the state of the file
 - Opened, closed, position etc.
- › `#include <stdio.h>`
 - contains many standard I/O functions and definitions for using files
- › **FILE** is a struct that is defined in `stdio.h` and this is the descriptor
- › To open a file, we use the `fopen` function

```
FILE *fopen(const char *path, const char *mode);
```

```
FILE * myfile = fopen("turtles.txt", "w");
```

variable

filename

mode

path can be relative, or absolute `/home/ssta7171/turtles.txt`

- › **FILE *fopen (...)**
 - modes
 - r** open text file for reading
 - w** truncate to zero length or create text file for writing
 - a** append; open or create text file for writing at end-of-file
 - rb** open binary file for reading
 - wb** truncate to zero length or create binary file for writing
 - ab** append; open or create binary file for writing at end-of-file
 - r+** open text file for update (reading and writing)
 - w+** truncate to zero length or create text file for update
 - a+** append; open or create text file for update, writing at end-of-file
 - › File versions of your lovable input/output
 - **fscanf**
 - **fprintf**
 - › Binary data
 - **fread**
 - **fwrite**
 - › Finish off with **fclose**
-

- › When your program begin, special files are opened for you:
 - `stdin`
 - `stdout`
 - `stderr`
- › You can use these files
 - `fscanf(stdin, ...)` same as `scanf(...)`
 - `fprintf(stdout, ...)` same as `printf(...)`
- › When a stream supports file position, the position is zero *
 - Every print/scan operation adjusts the position in the stream
 - Query position `ftell`, change position `fseek`

- › For reading input files, e.g. `stdin`, the end of file is important
 - `feof()` tests the end of file indicator
 - EOF does not happen until trying to read beyond end of stream

```
while ( ! feof(stdin) ) {  
    int num;  
    fscanf(stdin, "%d", &num);  
    fprintf(stderr, "num: %d\n", num);  
}
```

- › For reading input files, e.g. `stdin`, the end of file is important
 - `feof()` tests the end of file indicator
 - EOF does not happen until trying to read beyond end of stream

```
while ( ! feof(stdin) ) {  
    int num;  
    fscanf(stdin, "%d", &num);  
    fprintf(stderr, "num: %d\n", num);  
}
```

```
while ( ! feof(stdin) ) {  
    int num;  
    int nread = fscanf(stdin, "%d", &num);  
    if (nread <= 0)  
        break;  
    fprintf(stderr, "num: %d\n", num);  
}
```

- › unbuffered – input/output is passed on as soon as possible
 - › fully buffered – input/output is accumulated into a block then passed
 - › line buffered – the block size is based on the newline character
 - › Which do you get? Depends.
 - Device driver writers should consider `setvbuf` for optimal block size
 - › **fflush**
 - Output streams: force write all data,
 - Input streams: discard any unprocessed buffered data.
-

- › Many problems with `fscanf` with rules about whitespace, newlines or complex format string
 - › `fgets` reads **one line** of input and returning a string (with the newline character)
 - Use string processing functions to deal with the returned data
 - › Use `fgets` correctly, together with `feof` to distinguish read errors vs end of file.
 - it will make life easier
 - › `ferror` when you get that feeling...
-

```
#include <stdio.h>
#include <string.h>

#define BUFLLEN (64)

int main(int argc, char **argv) {
    int len;
    char buf[BUFLLEN];
    while (fgets(buf, BUFLLEN, stdin) != NULL) {
        len = strlen(buf);
        printf("%d\n", len);
    }
    return 0;
}
```
